



# Data Types & Feature Engineering



# Summary

- Introduction
- Structured data
- Semi-structured data
- Unstructured data
- Questions

# Data types

01



## Tabular

That is, a table is followed by the names of the columns, followed by a sequence of rows.

Example of data:

- numeric
- categorical
- boolean
- Timestamp
- Compound data (list, set, tuple, dict, serialize, ...)

02



## Image

Image classification  
(single-label)

Image classification  
(multi-label)

Image object detection

03



## Text

Single / Multi label  
classification

Entity recognition

Sentiment Analysis

Text translation

04



## Video

Action recognition

Classification


Object tracking

Video format:


- ❑ .MOV
- ❑ .MPEG4
- ❑ .MP4
- ❑ .AVI

# Feature Engineering

Feature engineering refers to the **process of using domain knowledge to select** and transform the most relevant variables from raw data when creating a predictive model using machine learning



**Data preparation** accounts for about **80%** of the work of data scientists



**What data scientists spend the most time doing**

- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets: 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Source: <https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/>

# Steps in Feature Engineering

01

## Feature Extraction

Feature extraction is the automatic creation of new variables by extracting them from raw data

The purpose of this step is to automatically reduce the volume of data into a more manageable set for modeling

02

## Data creation / preparation

Creating features involves identifying the variables that will be most useful in the predictive model

This is a subjective process that requires human intervention and creativity.

03

## Exploratory Data Analysis (EDA)

This step is used to identify and summarize the main characteristics in a data set through data analysis and investigation

Data scientist determine which statistical techniques are most appropriate for data analysis, and for choosing the right features for a model.

04

## Transformation / Selection

Transformation involves manipulating the predictor variables to improve model performance

Feature selection algorithms essentially analyze, judge, and rank various features to determine which features are irrelevant/redundant and should be removed

## Feature Engineering process

- Data preparation and exploratory data analysis
- Brainstorming/testing features and choosing which features to create
- Creating features
- Testing the impact of the identified features on the task;
- Optimizing the features if needed and repeating until the features work effectively

### Typical engineered features

The following list provides some typical ways to engineer useful features

- Numerical transformations (like scaling)
- Category encoder like one-hot encoder (for [categorical data](#))
- Clustering
- Group aggregated values
- Principal component analysis (PCA for numerical data)
- Data imputation handling outliers



# NLP

Natural Language Processing

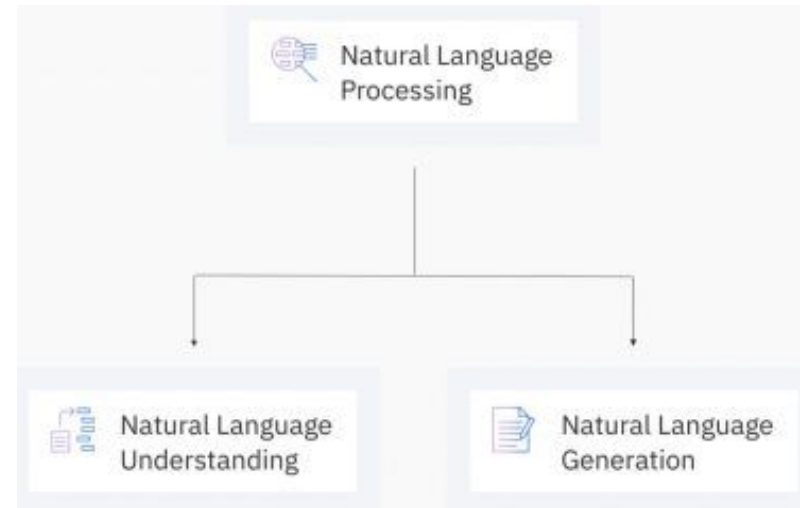


# What is NLP ?

**Natural language processing (NLP)** refers to the branch of computer science concerned with **giving computers the ability to understand text and spoken words in much the same way human beings can**. NLP seeks to convert unstructured language data into a structured data format to enable machines to understand speech and text and formulate relevant, contextual responses

**Natural language understanding (NLU)** uses **syntactic and semantic analysis of text and speech to determine the meaning** of a sentence. NLU focuses on machine reading comprehension through grammar and context, enabling it to determine the intended meaning of a sentence.

**Natural language generation (NLG)** is the process of **producing a human language text response based on some data input**. NLG enables computers to write. NLG focuses on text generation, or the construction of text in English or other languages, by a machine and based on a given dataset.



# NLP tasks

**Speech recognition**, also called speech-to-text, is the task of reliably converting voice data into text data.

**Named entity recognition**, or NEM, identifies words or phrases as useful entities. NEM identifies 'Kentucky' as a location or 'Fred' as a man's name.

**Part of speech tagging**, also called grammatical tagging, is the process of determining the part of speech of a particular word or piece of text based on its use and context. Part of speech identifies 'make' as a verb in 'I can make a paper plane,' and as a noun in 'What make of car do you own?'

**Sentiment analysis** attempts to extract subjective qualities-attitudes, emotions, sarcasm, confusion, suspicion-from text.

**Natural language generation (NLG)** is sometimes described as the opposite of speech recognition or speech-to-text; it's the task of putting structured information into human language.

**Word sense disambiguation** is the selection of the meaning of a word with multiple meanings through a process of semantic analysis that determine the word that makes the most sense in the given context. For example, word sense disambiguation helps distinguish the meaning of the verb 'make' in 'make the grade' (achieve) vs. 'make a bet' (place).

**Co-reference resolution** is the task of identifying if and when two words refer to the same entity. The most common example is determining the person or object to which a certain pronoun refers (e.g., 'she' = 'Mary'),

# NLP tools and approaches

**rules-based (RB) systems** could perform certain NLP tasks, but couldn't easily scale to accommodate a seemingly endless stream of exceptions or the increasing volumes of text and voice data.

**statistical NLP** combines computer algorithms with machine learning and deep learning models to automatically extract, classify, and label elements of text and voice data

## NLP use cases:

- Spam detection (gmail)


- Machine translation (Google translate, DeepL)

- Virtual agents and chatbots (Alexa, Siri, virtual assistant)

- Social media sentiment analysis (social media post, tweets, ecommerce reviews, ...)

- Text summarization (research databases, article keywords)

# ML algorithm processing only **numerical data**



## Terminologies

- **Document:** A document is a single text data point. For Example, a review of a particular product by the user.
- **Corpus:** It is a collection of all the documents present in our dataset.
- **Feature:** every unique token (word or sequence of words) in the corpus is considered as a feature.

## Binary vectorizer (bag of words)

good movie		<b>good</b>	<b>movie</b>	<b>not</b>	<b>a</b>	<b>did</b>	<b>like</b>
not a good movie	→	1	1	0	0	0	0
did not like		1	1	1	1	0	0
		0	0	1	0	1	1

It is obvious that it has some problems:

- We lose word order, hence the name “bag of words” as they're not ordered
- The counters are not normalized.

# Binary vectorizer (n-grams)

Example: 2-grams or bi-grams

- ...

good movie		<b>good movie</b>	<b>movie</b>	<b>did not</b>	<b>a</b>	<b>...</b>
not a good movie	→	<b>1</b>	<b>1</b>	<b>0</b>	<b>0</b>	<b>...</b>
did not like		<b>1</b>	<b>1</b>	<b>0</b>	<b>1</b>	<b>...</b>
		<b>0</b>	<b>0</b>	<b>1</b>	<b>0</b>	<b>...</b>

- The counters are not normalized
- Sparse matrix (contain mostly zero values)

# Term Frequency (TF) or Count vectorization

		text																			
0		Eddard Stark is a king in the north.																			
1		A king but one king : kings are everywhere.																			
2		Hodor was different : he was not a king .																			
3		But the North could not change without him.																			

	king	was	the	not	But	him	one	north	kings	is	in	he	Eddard	everywhere	different	could	change	but	are	Stark	North	Hodor	without
0	1	0	1	0	0	0	0	1	0	1	1	0	1	0	0	0	0	0	0	1	0	0	0
1	2	0	0	0	0	0	1	0	1	0	0	0	0	1	0	0	0	1	1	0	0	0	0
2	1	2	0	1	0	0	0	0	0	0	0	1	0	0	1	0	0	0	0	0	0	1	0
3	0	0	1	1	1	1	0	0	0	0	0	0	0	0	0	1	1	0	0	0	1	0	1

## Term frequency (TF)

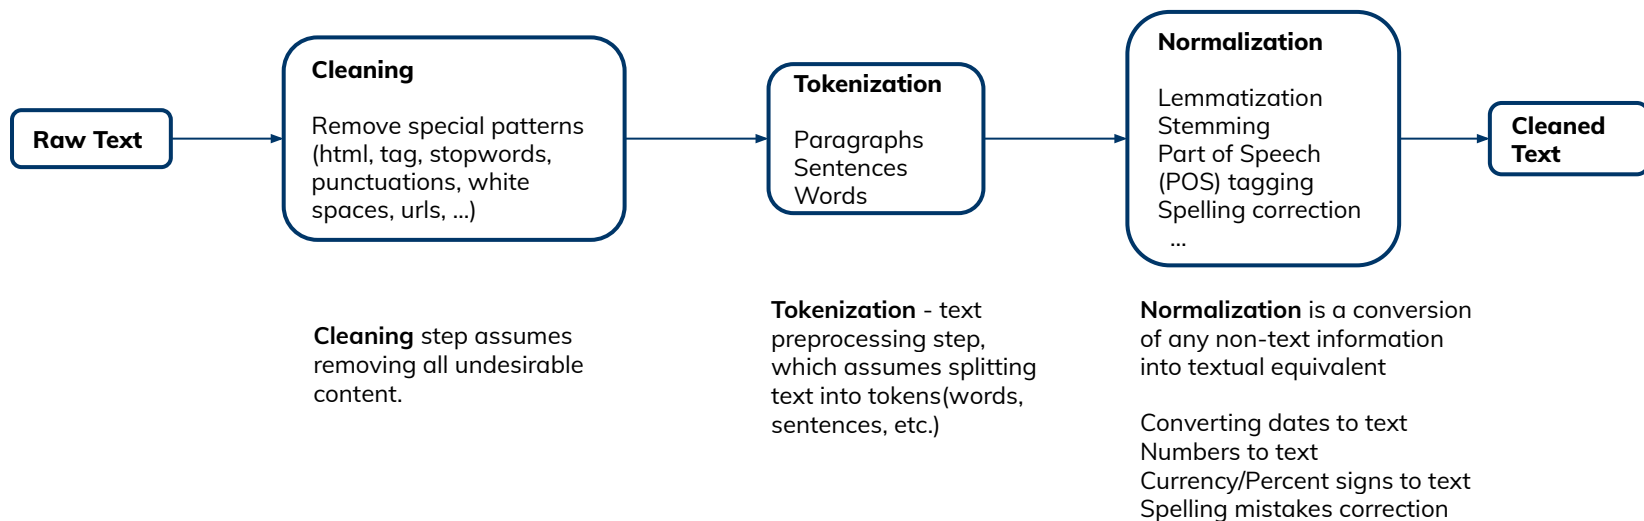
- $tf(t, d)$  – frequency for term (or n-gram)  $t$  in document  $d$
- Variants:

weighting scheme	TF weight
binary	0, 1
raw count	$f_{t,d}$
term frequency	$f_{t,d} / \sum_{t' \in d} f_{t',d}$
log normalization	$1 + \log(f_{t,d})$

## Drawbacks:

- The counters are not normalized
- Sparse matrix
- Terms are usually not statistically independent
- Missing semantic sensitivity

# What is NLP text preprocessing?



## Notes

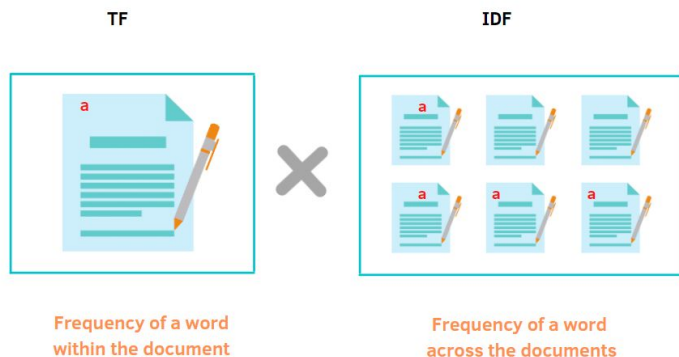
**Stemming** is the process of reducing inflection in words to their root forms such as mapping a group of words to the same stem even if the stem itself is not a valid word in the Language (ex: **studies** => **studi**)

**Lemmatization**, unlike Stemming, reduces the inflected words properly ensuring that the root word belongs to the language. Lemmatization considers the context and converts the word to its meaningful base form, which is called lemma. (ex: **studies** => **study**)



# TF-IDF

Term frequency - inverse document frequency



## Inverse document frequency (IDF)

- $N = |D|$  – total number of documents in corpus
- $|\{d \in D: t \in d\}|$  – number of documents where the term  $t$  appears
- $\text{idf}(t, D) = \log \frac{N}{|\{d \in D: t \in d\}|}$

## TF-IDF

- $\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$
- A high weight in TF-IDF is reached by a high term frequency (in the given document) and a low document frequency of the term in the whole collection of documents

Drawbacks:


- Sparse matrix, scalability (Very high dimensional)
- Low/High frequency
- Untreated synonym
- Terms are usually not statistically independent
- Missing semantic sensitivity

# Word embedding


**Word embedding** is a term used for the representation of words for text analysis, typically in the form of a real-valued vector that encodes the meaning of the word

Word embedding is able to capture the context, the semantic and syntactic similarity (gender, synonyms, ...) of a word by reducing the dimension. Word2Vec is the first neural embedding model


Methods to generate this mapping include [neural networks](#), [dimensionality reduction](#) on the word [co-occurrence matrix](#)



Male-Female



Verb tense

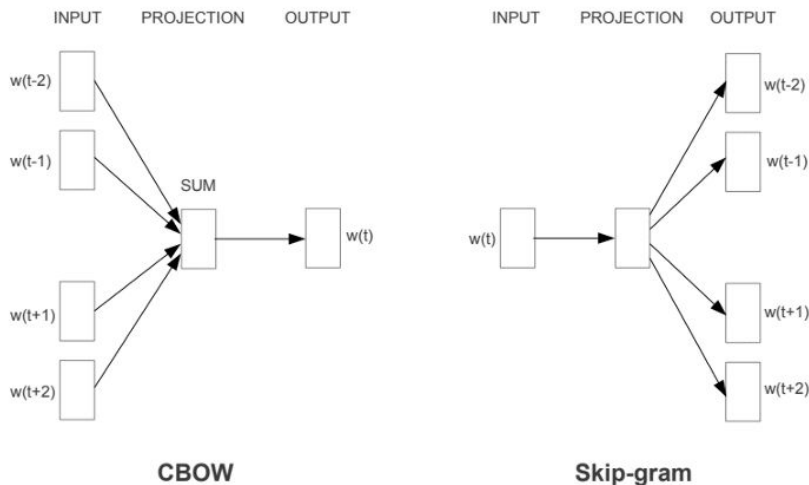


Country-Capital

# Word embedding

## Word2Vec

Word2Vec is a statistical method for efficiently learning a standalone word embedding from a text corpus. It was developed by Tomas Mikolov, et al. at Google in 2013



**Continuous Bowl of Words(CBOW)** try to fit the neighboring words in the window to the central word.

**Skip Gram** try to make the central word closer to the neighboring words. It is the complete opposite of the CBOW model

**Train your own word2vec model on a **custom** corpus**

**Pre-trained Word Embedding Models:**

- [Word2Vec](#) (Google)
- [GloVe](#) (Stanford, Global Vectors for Word Representation)
- [SpaCy](#)
- [fastText](#) (Facebook, Enriching Word Vectors with Subword Information)
- [Flair](#) (University of Berlin) etc.

# Word embedding

In very simple terms, Word Embeddings are the texts converted into numbers and there may be different numerical representations of the same text.

## Benefits of using Word Embeddings:

- It is much faster to train than hand build models like WordNet(which uses **graph embeddings**)
- Almost all modern NLP applications start with an embedding layer
- It Stores an approximation of meaning
- dense and low-dimensional vectors is computational

## Drawbacks of Word Embeddings:

- It can be memory intensive
- It is corpus dependent. Any underlying bias will have an effect on your model
- It cannot distinguish between homophones. Eg: brake/break, cell/sell, weather/whether etc.

# Word Embedding (CNN, RNN, ELMo, ...)

A proliferation of NLP programs, especially among the digital giants. Among the most advanced models, we can mention:

Google: **Bert, Albert**

## **Derived models**

- RoBERTa (Facebook)
- StructBERT (Alibaba)
- DeBERTa (Microsoft)
- DistilBERT (Hugging Face)

## **Alternative models**

- GPT-2 et GPT-3 (OpenAI)
- XLNet (Université Carnegie Mellon)
- UniLM (Microsoft)
- Reformer (Google)



# Image


Processing

# What is image processing?

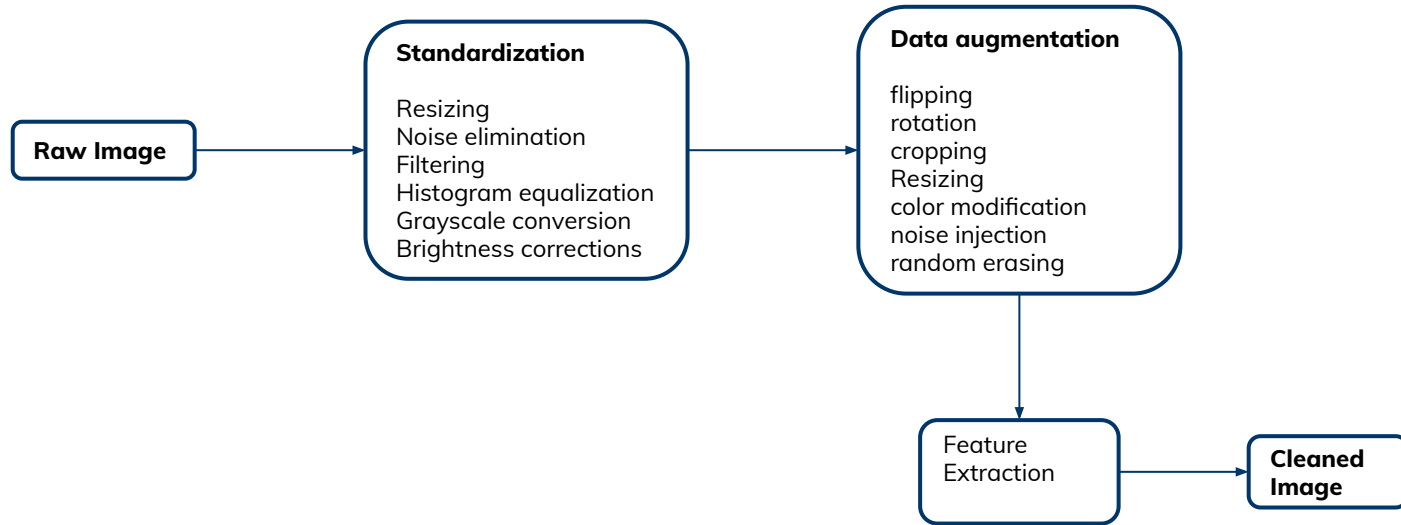
Image preprocessing is one of the most under-explored problems in the data science community. Image preprocessing converts image data into a form that allows machine learning algorithms to solve it. Some of the tools and platforms used in image preprocessing include: Pytorch, OpenCV, Keras, Tensorflow, and Pillow.

Example of use cases:

- Face detection
- Medical imaging
- Machine vision
- Autonomous Driving
- Image matching




# Image processing flowchart





# Image processing - Data augmentation



# Image processing - Feature Extraction

## Algorithms for Identification

- [Harris Corner](#)
- [SIFT\(Scale Invariant Feature Transform\)](#)
- [SURF\(Speeded Up Robust Feature\)](#)
- [FAST\(Features from Accelerated Segment Test\)](#)
- [ORB\(Oriented FAST and Rotated BRIEF\)](#)
- [BRIEF \(Binary Robust Independent Elementary Features\)](#)



# Best practices

# Best practices for creating training data

## → Avoid target leakage

Target leakage happens when your **training data includes predictive information that is not available when you ask for a prediction**. Target leakage can cause your model to show excellent evaluation metrics, but perform poorly on real data.

## → Avoid training-serving skew

Training-serving skew happens when you generate your **training data differently** than you generate the data you use to request **predictions**.

For example, if you use an average value, and for training purposes you average over 10 days, but when you request prediction, you average over the last month.

## → Avoid bias

Make sure that your **training data is representative of the entire universe** of potential data that you will be making predictions for. For example, if you have customers that live all over the world, you should not use training data from only one country.

## → Use extra care with imbalanced classes for classification models

If you have imbalanced classes (a classification problem with one or more outcomes that is seen rarely). Provide sufficient training data for the minority class or undersample the majority class (resampling method).

## → Provide enough training data

If you don't provide enough training data, the resulting model might perform poorly. The more columns you use to train your model, the more data you need to provide.

# References

<https://www.ibm.com/blogs/watson/2020/11/nlp-vs-nlu-vs-nlg-the-differences-between-three-natural-language-processing-concepts/>  
<https://www.ibm.com/cloud/learn/natural-language-processing>  
<https://www.geeksforgeeks.org/word-embeddings-in-nlp/>  
[https://fr.wikipedia.org/wiki/Word\\_embedding](https://fr.wikipedia.org/wiki/Word_embedding)  
<https://machinelearningmastery.com/what-are-word-embeddings/>  
<https://machinelearningmastery.com/develop-word-embeddings-python-gensim/>  
<https://medium.com/data-breach/introduction-to-orb-oriented-fast-and-rotated-brief-4220e8ec40cf>  
<https://prince-canuma.medium.com/image-pre-processing-c1aec0be3edf>  
<https://nanonets.com/blog/data-augmentation-how-to-use-deep-learning-when-you-have-limited-data-part-2/>  
<https://openclassrooms.com/fr/courses/4470531-classez-et-segmentez-des-donnees-visuelles/5024566-appliquez-vos-premiers-traitements-dimages>  
<https://www.omnisci.com/technical-glossary/feature-engineering>  
[https://en.wikipedia.org/wiki/Feature\\_engineering](https://en.wikipedia.org/wiki/Feature_engineering)  
<https://towardsdatascience.com/feature-engineering-for-machine-learning-3a5e293a5114>

**Questions ???**